

Step 1 | Data Pre-Processing

Importing the Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib
import seaborn as sns
sns.set(color_codes=True)
from sklearn.preprocessing import LabelEncoder

Importing the dataset

In [2]: dataset = pd.read_csv("Social_Network_Ads.csv")

In [3]: dataset.head()
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
In [4]: dataset.shape

Out[4]: (400, 5)

In [5]: dataset.columns

Out[5]: Index(['User ID', 'Gender', 'Age', 'EstimatedSalary', 'Purchased'], dtype='object')

In [6]: dataset.isna().any()

Out[6]: User ID      False
Gender      False
Age         False
EstimatedSalary  False
Purchased   False
dtype: bool

In [7]: dataset.isna().sum()

Out[7]: User ID      0
Gender      0
Age         0
EstimatedSalary  0
Purchased    0
dtype: int64

In [9]: dataset["Gender"].unique()

Out[9]: array(['Male', 'Female'], dtype=object)

In [10]: dataset["Gender"].replace({'Male':'0', 'Female':'1'},inplace=True)

In [11]: dataset
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	0	19	19000	0
1	15810944	0	35	20000	0
2	15668575	1	26	43000	0
3	15603246	1	27	57000	0
4	15804002	0	19	76000	0
...
395	15691863	1	46	41000	1
396	15706071	0	51	23000	1
397	15654296	1	50	20000	1
398	15755018	0	36	33000	0
399	15594041	1	49	36000	1

400 rows × 5 columns

Splitting the dataset into the Training set and Test set

```
In [13]: from sklearn.model_selection import train_test_split

In [14]: X_train, X_test, y_train, y_test = train_test_split(dataset[["Gender","Age","EstimatedSalary"]],dataset["Purchased"],test_size=0.2)

In [15]: len(X_train)

Out[15]: 320

In [16]: len(X_test)

Out[16]: 80

Feature Scaling

In [17]: from sklearn.preprocessing import StandardScaler

In [18]: sc_X = StandardScaler()

In [19]: X_train = sc_X.fit_transform(X_train)

In [20]: X_test = sc_X.transform(X_test)
```

Step 2 | Logistic Regression Model

```
In [21]: from sklearn.linear_model import LogisticRegression

In [22]: lr = LogisticRegression()
```

Fitting Logistic Regression to the Training set

```
In [23]: lr.fit(X_train,y_train)

Out[23]: LogisticRegression()
```

Step 3 | Predction

```
In [24]: lr.predict(X_test)

Out[24]: array([0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
        0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0], dtype=int64)

In [26]: X_test

Out[26]: array([[ 9.75304830e-01, -1.12010746e+00, -7.95427841e-01],
 [ -1.02532046e+00,  2.30291199e-01, -3.24934479e-01],
 [ -1.02532046e+00, -4.44908129e-01, -1.14829786e+00],
 [  9.75304830e-01, -2.51994035e-01,  2.26277902e+00],
 [ -1.02532046e+00,  7.12576434e-01, -1.11892020e+00],
 [  9.75304830e-01,  3.26748246e-01,  5.73413786e-02],
 [  9.75304830e-01, -5.90799412e-02,  1.96872066e+00],
 [  9.75304830e-01, -5.90799412e-02,  2.92588060e-01],
 [  9.75304830e-01,  1.58068986e+00,  1.11595144e+00],
 [  9.75304830e-01, -2.51994035e-01,  2.79355434e-02],
 [ -1.02532046e+00, -2.51994035e-01, -4.42557819e-01],
 [ -1.02532046e+00,  2.15943214e+00, -8.24839367e-01],
 [ -1.02532046e+00, -1.12010746e+00,  4.69023071e-01],
 [ -1.02532046e+00,  2.30291199e-01, -3.83746149e-01],
 [ -1.02532046e+00, -7.34279269e-01,  2.92588060e-01],
 [ -1.02532046e+00,  7.12576434e-01, -1.41295038e+00],
 [ -1.02532046e+00,  3.73771057e-02,  2.79355434e-02],
 [  9.75304830e-01,  1.96651804e+00,  7.33675587e-01],
 [ -1.02532046e+00, -1.21656450e+00,  2.63182225e-01],
 [  9.75304830e-01,  1.96651804e+00, -6.77804501e-01],
 [ -1.02532046e+00,  7.12576434e-01, -1.41295038e+00],
 [ -1.02532046e+00,  3.73771057e-02,  2.79355434e-02],
 [  9.75304830e-01,  1.96651804e+00,  7.33675587e-01],
 [ -1.02532046e+00, -1.21656450e+00,  2.63182225e-01],
 [  9.75304830e-01,  1.96651804e+00, -6.77804501e-01],
 [ -1.02532046e+00, -1.69884974e+00,  4.69023071e-01],
 [ -1.02532046e+00, -9.27193363e-01,  2.63182225e-01],
 [  9.75304830e-01, -5.41365175e-01,  1.90980899e+00],
 [ -1.02532046e+00, -1.40947860e+00, -2.07311138e-01],
 [ -1.02532046e+00,  4.23205293e-01, -1.77905303e-01],
 [ -1.02532046e+00, -2.51994035e-01, -9.13051182e-01],
 [ -1.02532046e+00, -5.11365175e-01, -1.53057372e+00],
 [ -1.02532046e+00, -1.55536988e-01,  8.51298928e-01],
 [ -1.02532046e+00, -7.34279269e-01, -1.5597956e+00],
 [ -1.02532046e+00,  4.23205293e-01, -4.71963654e-01],
 [  9.75304830e-01, -1.55536988e-01, -2.07311138e-01],
 [  9.75304830e-01,  1.09840462e+00,  2.08636401e+00],
 [  9.75304830e-01,  1.00194757e+00,  1.99812650e+00],
 [ -1.02532046e+00, -6.37822222e-01, -3.54340314e-01],
 [ -1.02532046e+00, -2.51994035e-01, -5.01369490e-01],
 [  9.75304830e-01,  2.30291199e-01,  1.45558884e-01],
 [  9.75304830e-01,  1.33834153e-01,  1.45558884e-01],
 [ -1.02532046e+00,  4.23205293e-01,  9.98328104e-01],
 [  9.75304830e-01, -2.51994035e-01, -1.38354454e+00],
 [ -1.02532046e+00, -9.27193363e-01,  5.57240576e-01],
 [ -1.02532046e+00,  1.33834153e-01,  1.88050316e+00],
 [ -1.02532046e+00, -2.51994035e-01,  8.67472137e-02],
 [  9.75304830e-01,  9.05490527e-01,  1.26298062e+00],
 [  9.75304830e-01,  3.26748246e-01,  2.92588060e-01],
 [ -1.02532046e+00, -1.12010746e+00, -1.17770370e+00],
 [ -1.02532046e+00, -6.37822222e-01,  1.16155049e-01],
 [ -1.02532046e+00, -5.90799412e-02, -5.30775325e-01],
 [  9.75304830e-01,  2.06297509e+00, -1.20710953e+00],
 [  9.75304830e-01,  1.96651804e+00, -9.42457017e-01],
 [  9.75304830e-01,  1.00194757e+00,  1.43941563e+00],
 [  9.75304830e-01, -7.34279269e-01,  1.35113813e+00],
 [  9.75304830e-01, -6.37822222e-01,  5.57240576e-01],
 [ -1.02532046e+00, -1.79530678e+00, -1.32473287e+00],
 [  9.75304830e-01,  9.05490527e-01, -5.89586995e-01],
 [ -1.02532046e+00, -1.12010746e+00, -1.58938539e+00],
 [  9.75304830e-01, -4.44908129e-01, -2.95528643e-01],
 [  9.75304830e-01,  1.96651804e+00, -1.38354454e+00],
 [ -1.02532046e+00, -1.55536988e-01, -2.95528643e-01],
 [  9.75304830e-01, -2.51994035e-01, -3.83746149e-01],
 [ -1.02532046e+00, -5.90799412e-02,  2.63182225e-01],
 [ -1.02532046e+00,  2.15943214e+00,  3.80805565e-01],
 [  9.75304830e-01,  2.06297509e+00,  1.74964719e-01],
 [ -1.02532046e+00, -1.55536988e-01, -1.08948619e+00],
 [  9.75304830e-01, -3.48451082e-01,  5.73413786e-02],
 [  9.75304830e-01, -1.55536988e-01,  1.45558884e-01],
 [  9.75304830e-01, -1.02365041e+00, -4.71963654e-01],
 [  9.75304830e-01, -6.37822222e-01, -1.61879123e+00],
 [  9.75304830e-01,  4.23205293e-01, -1.47029176e-03],
 [ -1.02532046e+00,  3.73771057e-02,  2.79355434e-02],
 [  9.75304830e-01,  1.67714690e+00,  1.61589054e+00],
 [  9.75304830e-01,  1.38717576e+00, -1.44235622e+00],
 [  9.75304830e-01, -1.79530678e+00, -1.29532704e+00],
 [ -1.02532046e+00, -8.30736316e-01,  2.29218485e+00],
 [  9.75304830e-01, -1.89176383e+00, -7.66022006e-01],
 [ -1.02532046e+00, -1.50593564e+00, -1.53057372e+00],
 [ -1.02532046e+00,  8.0903481e-01,  2.63182225e-01],
 [  9.75304830e-01,  1.00194757e+00, -1.03067452e+00],
 [  9.75304830e-01, -1.12010746e+00,  2.92588060e-01],
 [ -1.02532046e+00,  1.58068986e+00, -1.47029176e-03],
 [  9.75304830e-01, -1.69884974e+00,  3.51399730e-01]])
```

Step 4 | Evaluating The Predction

```
In [27]: lr.score(X_test,y_test)

Out[27]: 0.9
```

Making the Confusion Matrix

```
In [31]: # confusion matrix in sklearn
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

# actual values
actual = [0, 0, 0, 1, 0, 0, 0, 0, 1, 0]
# predicted values
predicted = [0, 0, 0, 1, 0, 0, 1, 0, 1]
# confusion matrix
matrix = confusion_matrix(actual,predicted, labels=[1,0])
print('Confusion matrix : \n',matrix)

# outcome values order in sklearn
tp, fn, fp, tn = confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : \n', tp, fn, fp, tn)

# classification report for precision, recall f1-score and accuracy
matrix = classification_report(actual,predicted,labels=[1,0])
print('Classification report : \n',matrix)

Confusion matrix :
[[1 1]
 [2 5]]
Outcome values :
1 1 2 5
Classification report :
              precision    recall  f1-score   support

      1         0.33         0.50         0.40         2
      0         0.83         0.71         0.77         7

   accuracy         0.58         0.61         0.58         9
  macro avg         0.72         0.67         0.69         9
 weighted avg         0.72         0.67         0.69         9
```

Visualization

